

Convergence of regularized agent-state based Q-learning in POMDPs

Amit Sinha (McGill), Matthieu Geist (Earth Species Project), Aditya Mahajan (McGill)

CDC 2025

- ▶ [email](mailto:aditya.mahajan@mcgill.ca): aditya.mahajan@mcgill.ca
- ▶ [web](https://adityam.github.io): <https://adityam.github.io>

Recent successes of RL

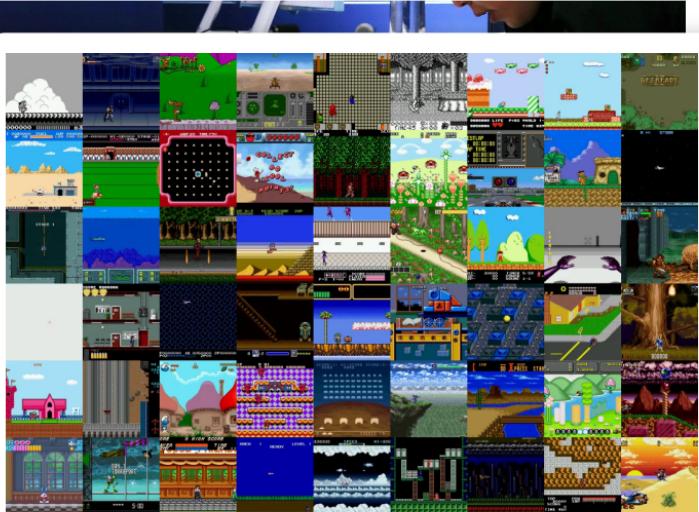
Recent successes of RL



Alpha Go

Regularized agent-state based Q-learning in POMDPs—(Sinha, Geist, Mahajan)

Recent successes of RL



Arcade games

Recent successes of RL



Robotic grasping

Recent successes of RL

- ▷ Algorithms based on comprehensive theory
- ▷ The theory is restricted almost exclusively to systems with **perfect state observations**



Robotic grasping

Recent successes of RL

- ▶ Algorithms based on comprehensive theory
- ▶ The theory is restricted almost exclusively to systems with **perfect state observations**

Many real-world applications are partially observed

- ▶ Healthcare
- ▶ Autonomous driving
- ▶ Finance (portfolio management)
- ▶ Retail and marketing



Robotic grasping

Recent successes of RL

- ▶ Algorithms based on comprehensive theory
- ▶ The theory is restricted almost exclusively to systems with **perfect state observations**



Robotic grasping

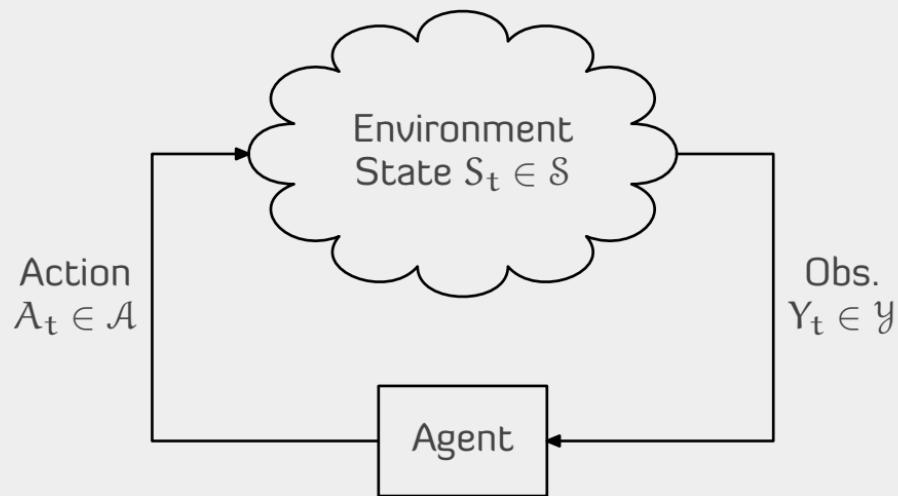
Many real-world applications are partially observed

- ▶ Healthcare
- ▶ Autonomous driving
- ▶ Finance (portfolio management)
- ▶ Retail and marketing

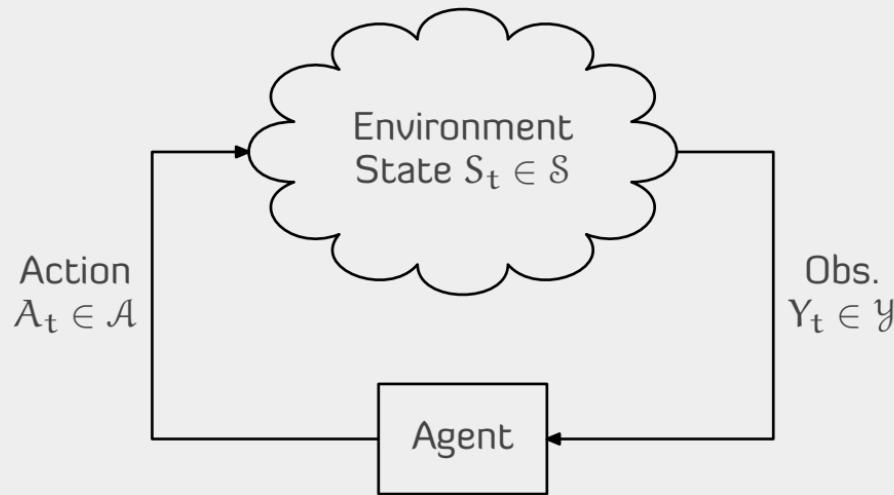
How do we develop a theory for RL for partially observed systems?

Regularized agent-state based Q-learning in POMDPs—(Sinha, Geist, Mahajan)

POMDPs: Partially observable Markov decision processes



POMDPs: Partially observable Markov decision processes



$$\begin{aligned}\mathbb{P}(S_{t+1}, Y_{t+1} | S_{1:t}, Y_{1:t}, A_{1:t}) \\ = \mathbb{P}(S_{t+1}, Y_{t+1} | S_t, A_t)\end{aligned}$$

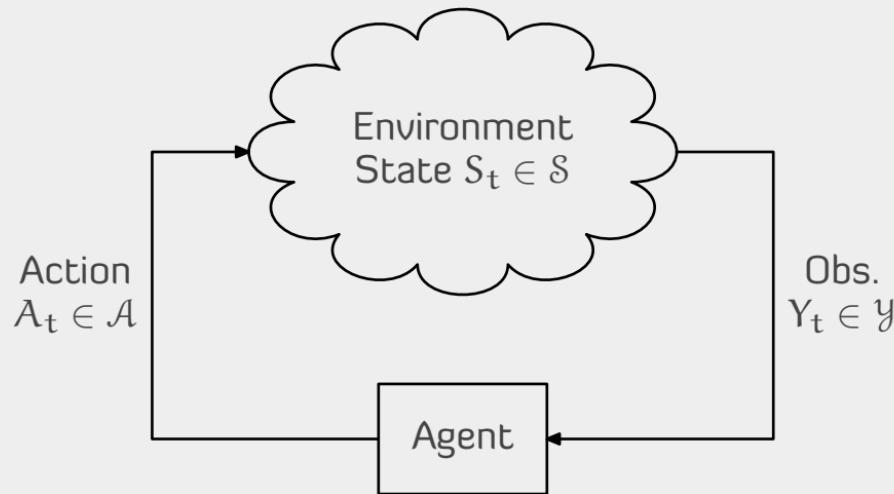
Reward: $R_t = r(S_t, A_t)$.

Policy: $\vec{\pi} = (\vec{\pi}_1, \vec{\pi}_2, \dots)$ where
 $A_t \sim \vec{\pi}_t(Y_{1:t}, A_{1:t-1})$

Performance:

$$J(\vec{\pi}) := \mathbb{E}^{\vec{\pi}} \left[\sum_{t=1}^{\infty} \gamma^{t-1} R_t \mid S_1 \sim \xi_1 \right]$$

POMDPs: Partially observable Markov decision processes



$$\begin{aligned}\mathbb{P}(S_{t+1}, Y_{t+1} | S_{1:t}, Y_{1:t}, A_{1:t}) \\ = \mathbb{P}(S_{t+1}, Y_{t+1} | S_t, A_t)\end{aligned}$$

Reward: $R_t = r(S_t, A_t)$.

Policy: $\vec{\pi} = (\vec{\pi}_1, \vec{\pi}_2, \dots)$ where
 $A_t \sim \vec{\pi}_t(Y_{1:t}, A_{1:t-1})$

Performance:

$$J(\vec{\pi}) := \mathbb{E}^{\vec{\pi}} \left[\sum_{t=1}^{\infty} \gamma^{t-1} R_t \mid S_1 \sim \xi_1 \right]$$

Objective: Find the (history-dependent) policy $\vec{\pi}$ that maximizes $J(\vec{\pi})$

Review: Belief-state based planning

Key simplifying idea

Define **belief state** $B_t \in \Delta(\mathcal{S})$ as $B_t(s) = \mathbb{P}(S_t = s | Y_{1:t}, A_{1:t-1})$.

- ▶ Belief state updates in a state-like manner: $B_{t+1} = \text{function}(B_t, Y_{t+1}, A_t)$.
- ▶ Belief state is sufficient to evaluate rewards: $\mathbb{E}[R_t | Y_{1:t}, A_{1:t}] = \hat{r}(B_t, A_t)$.

Thus, $\{B_t\}_{t \geq 1}$ is a **perfectly observed** controlled Markov process.

- ❑ Astrom, "Optimal control of Markov processes with incomplete information," JMAA 1965.
- ❑ Stratonovich, "Conditional Markov Processes," TVP 1960.

Review: Belief-state based planning

Key simplifying idea

Define **belief state** $B_t \in \Delta(\mathcal{S})$ as $B_t(s) = \mathbb{P}(S_t = s | Y_{1:t}, A_{1:t-1})$.

- ▶ Belief state updates in a state-like manner: $B_{t+1} = \text{function}(B_t, Y_{t+1}, A_t)$.
- ▶ Belief state is sufficient to evaluate rewards: $\mathbb{E}[R_t | Y_{1:t}, A_{1:t}] = \hat{r}(B_t, A_t)$.

Thus, $\{B_t\}_{t \geq 1}$ is a **perfectly observed** controlled Markov process. Therefore:

Structure of optimal policy

There is no loss of optimality in choosing the action A_t as a function of the belief state B_t

Dynamic Program

The optimal control policy is given a DP with belief B_t as state.

Implications of the POMDP modeling framework

Implications for planning

- ▶ Allows the use of the MDP machinery for partially observed sys.
- ▶ Various exact and approximate algorithms to efficiently solve DP.
Exact: incremental pruning, witness algorithm, linear support algo
Approximate: QMDP, point based methods, SARSOP, DESPOT,

Implications of the POMDP modeling framework

- ▶ Allows the use of the MDP machinery for partially observed sys.
- ▶ Allows the use of standard RL algorithms
- ▶ The construction of the belief state depends on the system model.
- ▶ So, when the system model is unknown, we cannot construct the belief state and therefore cannot use standard RL algorithms.

Implications for learning

Implications of the POMDP modeling framework

Implications for learning

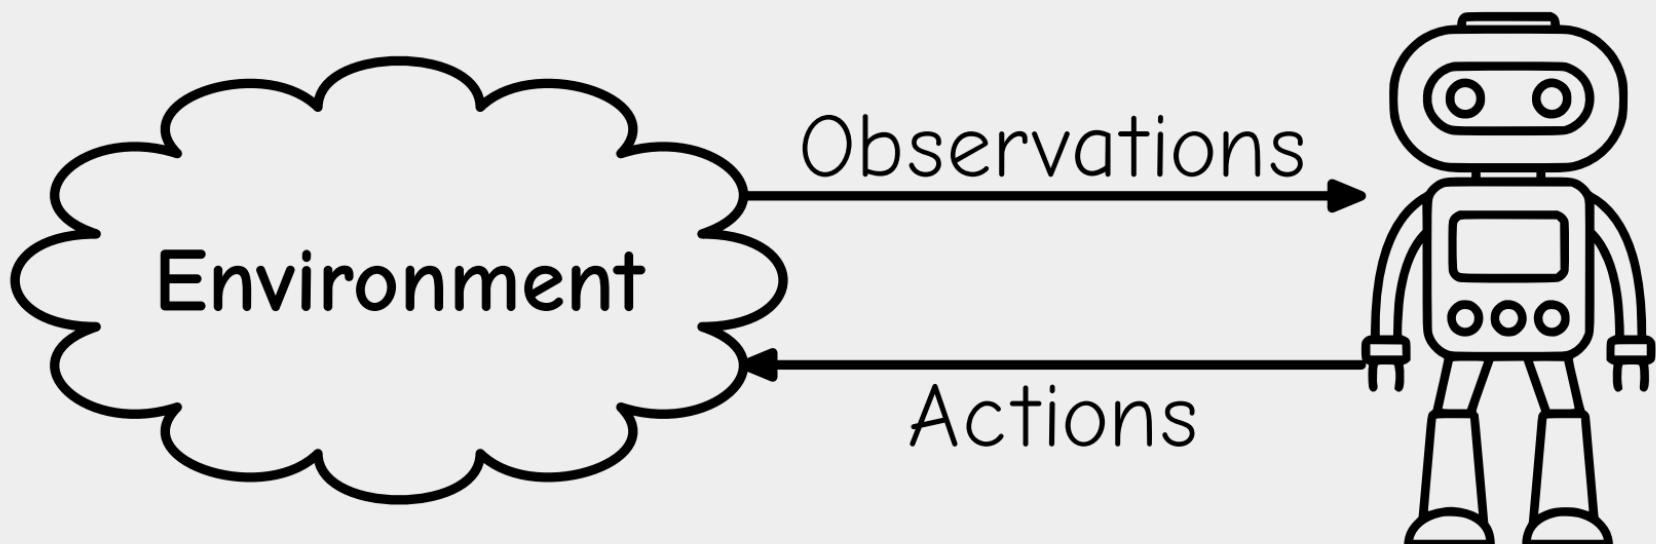
- ▷ Allows the use of the MDP machinery for partially observed sys.
- ▷ Many algorithms have been proposed for POMDP
- ▷ The construction of the belief state depends on the system model.
- ▷ So, when the system model is unknown, we cannot construct the belief state and therefore cannot use standard RL algorithms.
- ▷ **On the theoretical side:**
 - ▷ Propose alternative methods: PSRs (predictive state representations), bisimulation metrics, . . .
 - ▷ Good theoretical guarantees, but difficult to scale.

Implications of the POMDP modeling framework

Implications for learning

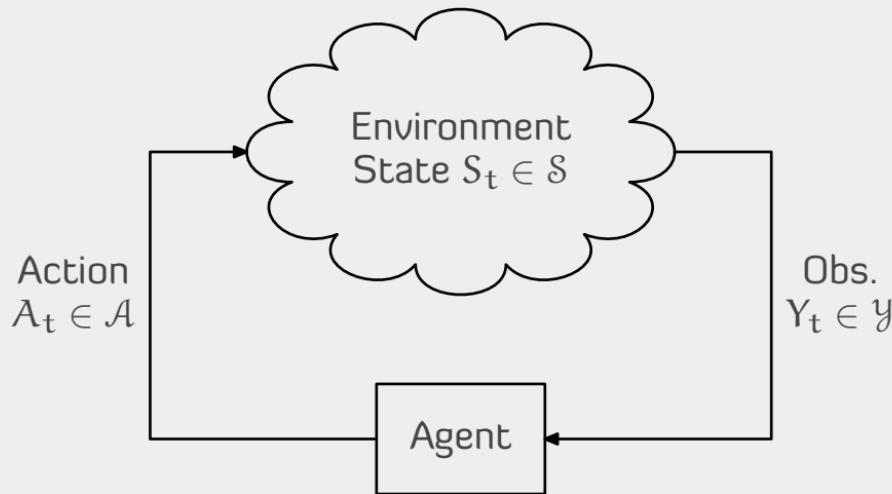
- ▷ Allows the use of the MDP machinery for partially observed sys.
- ▷ Various ways to handle the unknown system model.
- ▷ The construction of the belief state depends on the system model.
- ▷ So, when the system model is unknown, we cannot construct the belief state and therefore cannot use standard RL algorithms.
- ▷ **On the theoretical side:**
 - ▷ Propose alternative methods: PSRs (predictive state representations), bisimulation metrics, . . .
 - ▷ Good theoretical guarantees, but difficult to scale.
- ▷ **On the practical side:**
 - ▷ Simply stack the previous k observations and treat it as a “state”.
 - ▷ Instead of a CNN, use an RNN to model policy and action-value fn.
 - ▷ Can be made to work but lose theoretical guarantees and insights.

Deep RL learns agent-state based policies



Can we understand the
convergence of such algorithms?

Abstract model of agent-state based policies



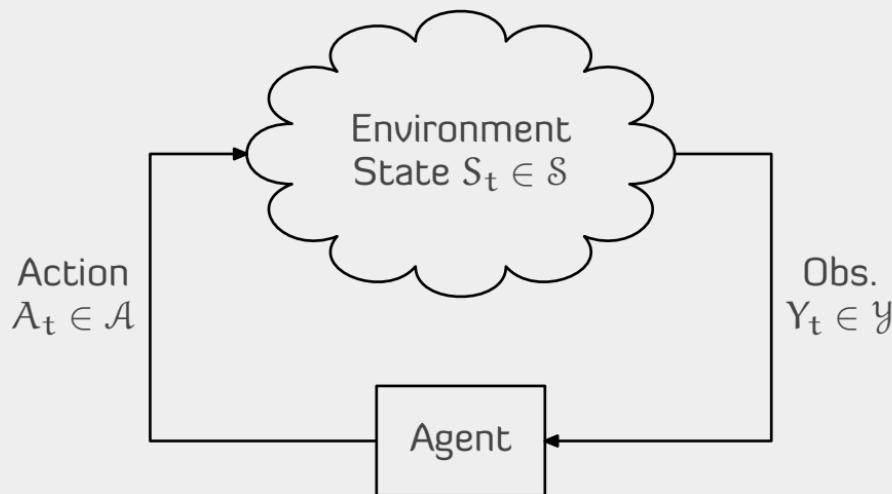
Agent state: $Z_t \in \mathcal{Z}$, where

$$Z_{t+1} = \phi(Z_t, Y_{t+1}, A_t)$$

Examples:

- ▷ $Z_t = (Y_{t-n:t}, A_{t-n:t-1})$
- ▷ Finite-state controllers
- ▷ Recurrent neural networks

Abstract model of agent-state based policies



Agent state: $Z_t \in \mathcal{Z}$, where

$$Z_{t+1} = \phi(Z_t, Y_{t+1}, A_t)$$

Examples:

- ▷ $Z_t = (Y_{t-n:t}, A_{t-n:t-1})$
- ▷ Finite-state controllers
- ▷ Recurrent neural networks

Regularized agent-state based Q-learning (RASQL)

- ▷ Act according to behavior policy $\mu: \mathcal{Z} \rightarrow \Delta(\mathcal{A})$
- ▷ Recursively update:

$$Q_{t+1}(z, a) = Q_t(z, a) + \alpha_t(z, a) [r_t + \gamma \Omega^*(Q_t(z_{t+1}, \cdot)) - Q_t(z, a)]$$

where Ω^* is **reward regularization** (replaces standard max operator).

Why regularization?

All empirical algorithms use regularization in some form.

- ▷ to encourage exploration
- ▷ to stabilize the learning process by smoothing the optimization landscape
- ▷ to constrain policy update to stay similar to previous policy

Why regularization?

All empirical algorithms use regularization in some form.

- ▷ to encourage exploration
- ▷ to stabilize the learning process by smoothing the optimization landscape
- ▷ to constrain policy update to stay similar to previous policy

Existing analysis of QL for POMDPs

Restricted to unregularized QL

- ▷ QL with finite window memory [Jaakkola Singh Jordan 1994, Kara Yuksel 2022]
- ▷ periodic QL with agent state [Sinha Geist Mahajan 2024]
- ▷ QL for non-Markovian environments [Kara Yuksel 2024, Chandak Shah Borkar 2024]

This paper: Analyze convergence
of regularized agent-state based
Q-learning (RASQL) for POMDPs

Key Tool: Legendre-Fenchel transform

Convex conjugate

For a strongly convex function $\Omega: \mathbb{R}^n \rightarrow \mathbb{R}$, its convex conjugate $\Omega^*: \mathbb{R}^n \rightarrow \mathbb{R}$ is defined as:

$$\Omega^*(q) = \max_{p \in \mathbb{R}^n} \{\langle p, q \rangle - \Omega(p)\}.$$

- Key tool for analyzing convergence of regularized MDPs [Geist et al 2019]

Key Tool: Legendre-Fenchel transform

Convex conjugate

For a strongly convex function $\Omega: \mathbb{R}^n \rightarrow \mathbb{R}$, its convex conjugate $\Omega^*: \mathbb{R}^n \rightarrow \mathbb{R}$ is defined as:

$$\Omega^*(q) = \max_{p \in \mathbb{R}^n} \{ \langle p, q \rangle - \Omega(p) \}.$$

- Key tool for analyzing convergence of regularized MDPs [Geist et al 2019]

Example: Entropy regularization

- $\Omega(p) = \frac{1}{\beta} \sum p(a) \ln p(a)$
- $\Omega^*(q) = \frac{1}{\beta} \ln \left(\sum \exp(\beta q(a)) \right)$

Key Tool: Legendre-Fenchel transform

Convex conjugate

For a strongly convex function $\Omega: \mathbb{R}^n \rightarrow \mathbb{R}$, its convex conjugate $\Omega^*: \mathbb{R}^n \rightarrow \mathbb{R}$ is defined as:

$$\Omega^*(q) = \max_{p \in \mathbb{R}^n} \{ \langle p, q \rangle - \Omega(p) \}.$$

- Key tool for analyzing convergence of regularized MDPs [Geist et al 2019]

Example: Entropy regularization

- $\Omega(p) = \frac{1}{\beta} \sum p(a) \ln p(a)$
- $\Omega^*(q) = \frac{1}{\beta} \ln \left(\sum \exp(\beta q(a)) \right)$

Example: KL regularization

- $\Omega(p) = \frac{1}{\beta} \sum p(a) \ln \frac{p(a)}{p_{\text{ref}}(a)}$
- $\Omega^*(q) = \frac{1}{\beta} \ln \left(\sum p_{\text{ref}}(a) \exp(\beta q(a)) \right)$

Key Tool: Legendre-Fenchel transform

Convex conjugate

For a strongly convex function $\Omega: \mathbb{R}^n \rightarrow \mathbb{R}$, its convex conjugate $\Omega^*: \mathbb{R}^n \rightarrow \mathbb{R}$ is defined as:

$$\Omega^*(q) = \max_{p \in \mathbb{R}^n} \{ \langle p, q \rangle - \Omega(p) \}.$$

- Key tool for analyzing convergence of regularized MDPs [Geist et al 2019]

Example: Entropy regularization

- $\Omega(p) = \frac{1}{\beta} \sum p(a) \ln p(a)$
- $\Omega^*(q) = \frac{1}{\beta} \ln \left(\sum \exp(\beta q(a)) \right)$

- **Important property:** The optimal regularized policy p^* (the argmax) can be obtained directly from the gradient of the convex conjugate:

$$p^* = \nabla \Omega^*(q)$$

Example: KL regularization

- $\Omega(p) = \frac{1}{\beta} \sum p(a) \ln \frac{p(a)}{p_{\text{ref}}(a)}$
- $\Omega^*(q) = \frac{1}{\beta} \ln \left(\sum p_{\text{ref}}(a) \exp(\beta q(a)) \right)$

Main step: construction of a limit MDP

Assumptions:

- The behavior policy μ induces an invariant distribution ζ_μ over (S_t, Y_t, Z_t, A_t) .
- The learning rates $\{\alpha_t\}_{t \geq 1}$ satisfy: $\sum_{t=1}^{\infty} \alpha_t = \infty$ and $\sum_{t=1}^{\infty} \alpha_t^2 < \infty$.

Main step: construction of a limit MDP

Assumptions:

- ▶ The behavior policy μ induces an invariant distribution ζ_μ over (S_t, Y_t, Z_t, A_t) .
- ▶ The learning rates $\{\alpha_t\}_{t \geq 1}$ satisfy: $\sum_{t=1}^{\infty} \alpha_t = \infty$ and $\sum_{t=1}^{\infty} \alpha_t^2 < \infty$.

Limit MDP

Construct a limit MDP with state space \mathcal{Z} and action space \mathcal{A} and

- ▶ Reward: $r_\mu(z, a) := \sum_{s \in \mathcal{S}} r(s, a) \zeta_\mu(s | z)$
- ▶ Dynamics: $P_\mu(z' | z, a) := \sum_{(s, y')} \mathbb{1}_{\{z' = \phi(z, y', a)\}} P(y' | s, a) \zeta_\mu(s | z)$

Main step: construction of a limit MDP

Assumptions:

- ▶ The behavior policy μ induces an invariant distribution ζ_μ over (S_t, Y_t, Z_t, A_t) .
- ▶ The learning rates $\{\alpha_t\}_{t \geq 1}$ satisfy: $\sum_{t=1}^{\infty} \alpha_t = \infty$ and $\sum_{t=1}^{\infty} \alpha_t^2 < \infty$.

Limit MDP

Construct a limit MDP with state space \mathcal{Z} and action space \mathcal{A} and

- ▶ Reward: $r_\mu(z, a) := \sum_{s \in \mathcal{S}} r(s, a) \zeta_\mu(s | z)$
- ▶ Dynamics: $P_\mu(z' | z, a) := \sum_{(s, y')} \mathbb{1}_{\{z' = \phi(z, y', a)\}} P(y' | s, a) \zeta_\mu(s | z)$
- ▶ Let Q_μ be the unique fixed point of the regularized Bellman operator of this limit MDP:

$$Q_\mu(z, a) := r_\mu(z, a) + \gamma \sum_{z' \in \mathcal{Z}} P_\mu(z' | z, a) \Omega^*(Q_\mu(z', \cdot))$$

Main convergence result

Theorem

Under the stated assumptions, RASQL converges to Q_μ almost surely.

Main convergence result

Theorem

Under the stated assumptions, RASQL converges to Q_μ almost surely.

Salient features

- ▷ The converged policy is the greedy regularized policy with respect to Q_μ :
$$\pi^*(\cdot | z) = \nabla \Omega^*(Q_\mu(z, \cdot))$$
- ▷ This policy is typically **stochastic** (e.g., softmax for entropy regularization)
- ▷ **Significant advantage over (unregularized) ASQL**, which converges to deterministic policy as stochastic stationary policies can outperform deterministic stationary policies in POMDPs with agent-states.

Generalization: Regularized periodic Q-learning (RePASQL)

Why periodic policies?

- ▶ Time-varying policies can outperform time-invariant ones (as agent state is not info state)
- ▶ Periodic policies are simplest time-varying policies.

Generalization: Regularized periodic Q-learning (RePASQL)

Why periodic policies?

- ▶ Time-varying policies can outperform time-invariant ones (as agent state is not info state)
- ▶ Periodic policies are simplest time-varying policies.

Regularized periodic Q-learning (RePASQL)

Fix a period L and pick a periodic behavior policy μ with period L . Recursively update:

$$Q_{t+1}^\ell(z, a) = Q_t^\ell(z, a) + \alpha_t^\ell(z, a) \left[r_t + \gamma \Omega^*(Q_t^{[\ell+1]}(z', \cdot)) - Q_t^\ell(z, a) \right]$$

where $[\ell] = \ell \bmod L$.

Generalization: Regularized periodic Q-learning (RePASQL)

Why periodic policies?

- ▶ Time-varying policies can outperform time-invariant ones (as agent state is not info state)
- ▶ Periodic policies are simplest time-varying policies.

Regularized periodic Q-learning (RePASQL)

Fix a period L and pick a periodic behavior policy μ with period L . Recursively update:

$$Q_{t+1}^\ell(z, a) = Q_t^\ell(z, a) + \alpha_t^\ell(z, a) \left[r_t + \gamma \Omega^*(Q_t^{[\ell+1]}(z', \cdot)) - Q_t^\ell(z, a) \right]$$

where $[\ell] = \ell \bmod L$.

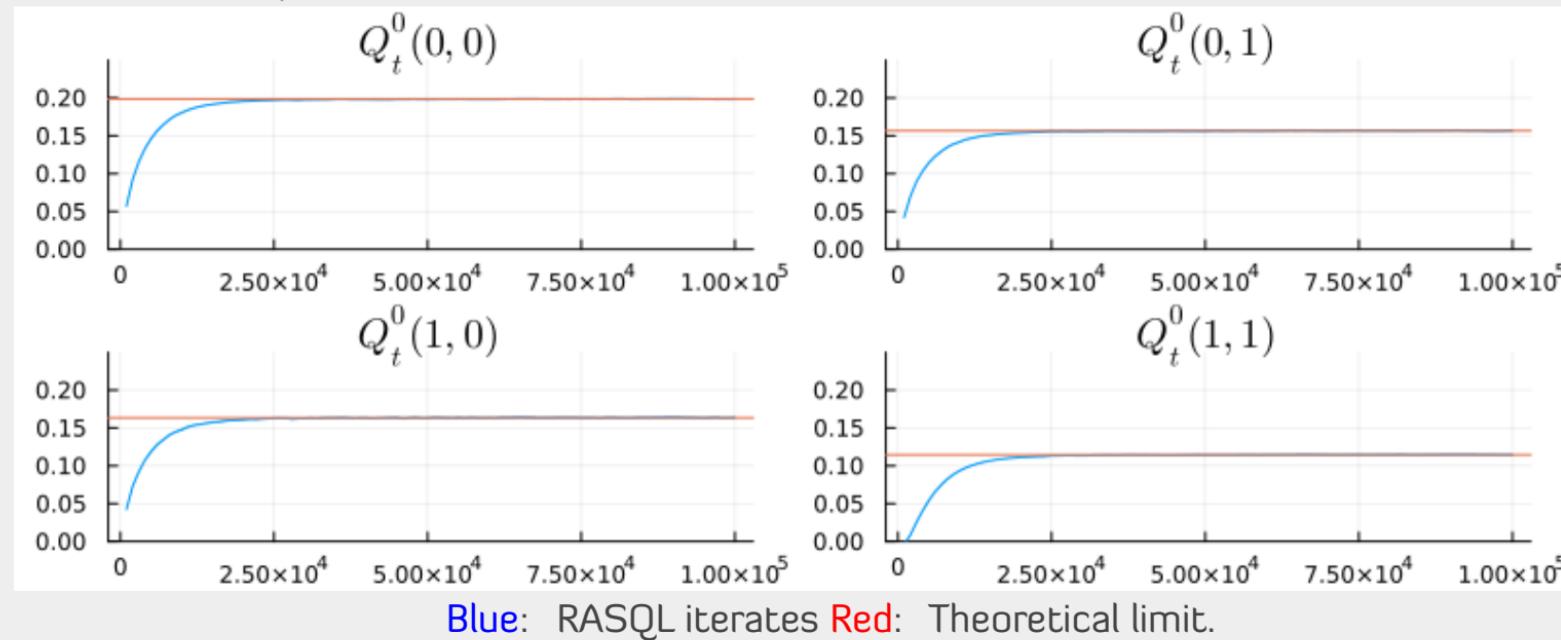
Theorem

Assuming (S_t, Y_t, Z_t, A_t) has periodic limit, then RePASQL converges to the fixed point of a periodic regularized MDP almost surely.

Numerical experiments

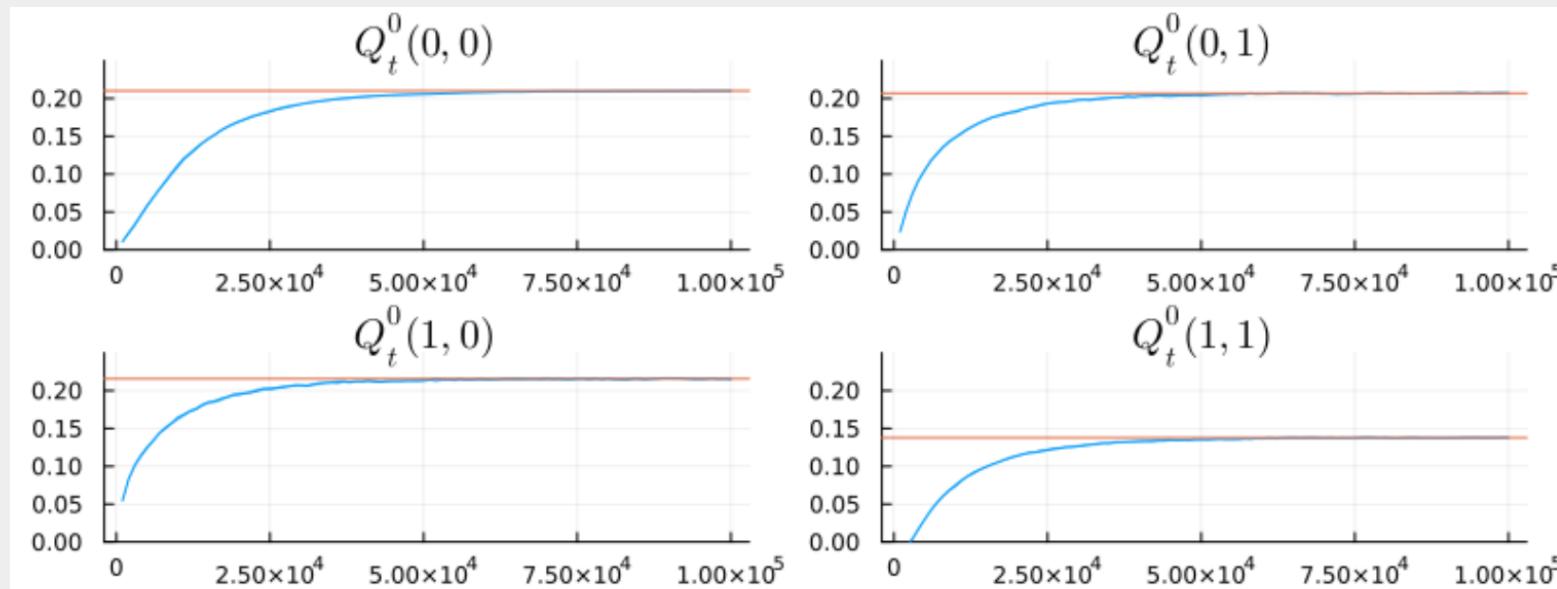
Verifying convergence of RASQL

- ▶ Small POMDP with $|\mathcal{S}| = 4$, $|\mathcal{Y}| = 2$, $|\mathcal{A}| = 2$
- ▶ Agent state $Z_t = Y_t$
- ▶ Median and quantile over 25 seeds



Verifying convergence of RePASQL

- ▶ Same model as before
- ▶ Period $L = 2$
- ▶ $\ell = 0$

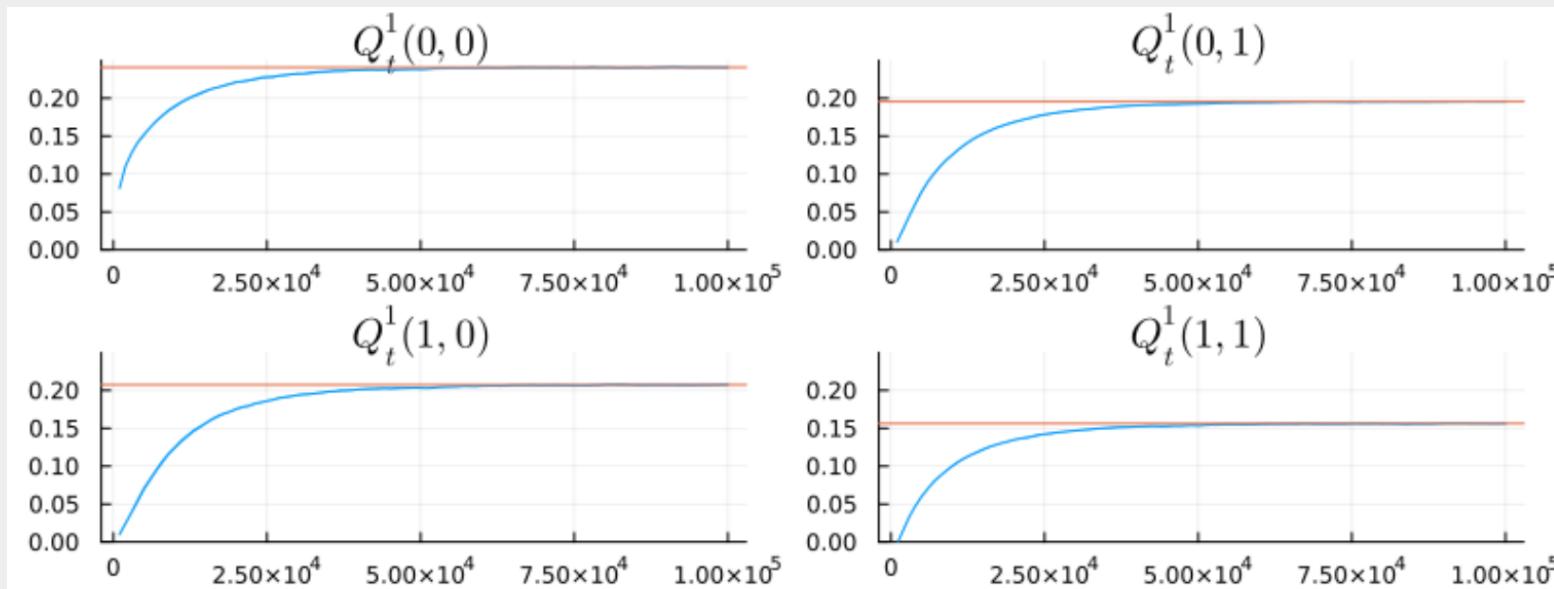


Blue: RePASQL iterates Red: Theoretical limit.

Regularized agent-state based Q-learning in POMDPs—(Sinha, Geist, Mahajan)

Verifying convergence of RePASQL

- ▶ Same model as before
- ▶ Period L = 2
- ▶ $\ell = 1$



Blue: RePASQL iterates Red: Theoretical limit.

Regularized agent-state based Q-learning in POMDPs—(Sinha, Geist, Mahajan)

Conclusion

- ▶ Proved almost sure convergence of **regularized** agent-state based Q-learning
- ▶ Iterates converge to the fixed point of a limit MDP (or a periodic limit MDP), which depend on the invariant distribution (or periodic limit distribution) induced by the exploration policy.
- ▶ Regularization helps in learning stochastic policies, which can outperform deterministic ones in POMDPs with agent-states.

Conclusion

- ▶ Proved almost sure convergence of **regularized** agent-state based Q-learning
- ▶ Iterates converge to the fixed point of a limit MDP (or a periodic limit MDP), which depend on the invariant distribution (or periodic limit distribution) induced by the exploration policy.
- ▶ Regularization helps in learning stochastic policies, which can outperform deterministic ones in POMDPs with agent-states.

- ▷ email: aditya.mahajan@mcgill.ca
- ▷ web: <https://adityam.github.io>

Thank you