

Learning to control networked-coupled subsystems with unknown dynamics

Aditya Mahajan
McGill University

Symcomore group meeting
29 Sep 2023

- ▶ email: aditya.mahajan@mcgill.ca
- ▶ homepage: <http://cim.mcgill.ca/~adityam>


Networks are ubiquitous

Networks are ubiquitous




Transportation networks

Networks are ubiquitous



Internet of Things

Networks are ubiquitous




Energy network

Networks are ubiquitous


Salient Features

- ▶ Large/growing size
- ▶ Nodes have local states
- ▶ Coupled dynamics and costs



Energy network

Networks are ubiquitous



Energy network


Salient Features

- ▶ Large/growing size
- ▶ Nodes have local states
- ▶ Coupled dynamics and costs

Design challenges

- ▶ Scalability of the solution
- ▶ How to handle model uncertainty

Networks are ubiquitous



Salient Features

- ▶ Large/growing size
- ▶ Nodes have local states
- ▶ Coupled dynamics and costs

Design challenges

- ▶ Scalability of the solution
- ▶ How to handle model uncertainty

Present a spectral decomposition method for network-coupled subsystems which leads to scalable planning and learning

Outline



System Model


- ▶ Network-coupled subsystems
 - ▶ Agents interacting over a graph
 - ▶ Coupled dynamics
 - ▶ Coupled costs

Outline



System Model

- ▶ Network-coupled subsystems
 - ▶ Agents interacting over a graph
 - ▶ Coupled dynamics
 - ▶ Coupled costs



Planning solution


- ▶ Spectral factorization of dynamics and cost
- ▶ Decoupled Riccati equations

Outline




System Model

- ▶ Network-coupled subsystems
 - ▶ Agents interacting over a graph
 - ▶ Coupled dynamics
 - ▶ Coupled costs



Planning solution

- ▶ Spectral factorization
of dynamics and cost
- ▶ Decoupled Riccati equations



Learning solution


- ▶ Spectral factorization of learning
- ▶ Numerical experiments

Outline




System Model

- ▶ Network-coupled subsystems
- ▶ Agents interacting over a graph
- ▶ Coupled dynamics
- ▶ Coupled costs



Planning solution

- ▶ Spectral factorization of dynamics and cost
- ▶ Decoupled Riccati equations




Learning solution

- ▶ Spectral factorization of learning
- ▶ Numerical experiments

System Model

Weighted undirected graph \mathcal{G}


- ▶ Nodes $N = \{1, \dots, n\}$.
- ▶ Symmetric matrix $M = [m^{ij}]$ associated with \mathcal{G} (e.g., weighted adjacency matrix, weighted Laplacian, etc.)



System Model

Weighted undirected graph \mathcal{G}

- ▶ Nodes $N = \{1, \dots, n\}$.
- ▶ Symmetric matrix $M = [m^{ij}]$ associated with \mathcal{G} (e.g., weighted adjacency matrix, weighted Laplacian, etc.)



System Dynamics


- ▶ A subsystem located at each node. **State** $x_t^i \in \mathbb{R}^{d_x}$. **Control** $u_t^i \in \mathbb{R}^{d_u}$.

$$x_{t+1}^i = Ax_t^i + Bu_t^i + D \sum_{j \in N} m^{ij} x_t^j + E \sum_{j \in N} m^{ij} u_t^j + w_t^i$$

System Model

Weighted undirected graph \mathcal{G}

- ▶ Nodes $N = \{1, \dots, n\}$.
- ▶ Symmetric matrix $M = [m^{ij}]$ associated with \mathcal{G} (e.g., weighted adjacency matrix, weighted Laplacian, etc.)



System Dynamics

- ▶ A subsystem located at each node. **State** $x_t^i \in \mathbb{R}^{d_x}$. **Control** $u_t^i \in \mathbb{R}^{d_u}$.

$$x_{t+1}^i = Ax_t^i + Bu_t^i + D \sum_{j \in N} m^{ij} x_t^j + E \sum_{j \in N} m^{ij} u_t^j + w_t^i$$

Network field of states $x_t^{\mathcal{G},i}$


Network field of control $u_t^{\mathcal{G},i}$

System Model (cont.)

Per-step cost

$$c(x_t, u_t) = \sum_{i,j \in \mathcal{N}} [\mathbf{h}_q^{ij}(x_t^i)^\top Q(x_t^j) + \mathbf{h}_r^{ij}(u_t^i)^\top Q(u_t^j)]$$

where $H_q = [h_q^{ij}]$ and $H_r = [h_r^{ij}]$ are **symmetric matrices** which have the same eigenvectors as M .



System Model (cont.)

Per-step cost


$$c(x_t, u_t) = \sum_{i,j \in \mathcal{N}} [\mathbf{h}_q^{ij}(x_t^i)^\top Q(x_t^j) + \mathbf{h}_r^{ij}(u_t^i)^\top Q(u_t^j)]$$

where $H_q = [h_q^{ij}]$ and $H_r = [h_r^{ij}]$ are **symmetric matrices which have the same eigenvectors as M** .

Remark

For two symmetric $n \times n$ matrices M_1 and M_2 , the following statements are **equivalent**:

- ▶ M_1 and M_2 share the same eigenvectors.
- ▶ M_1 and M_2 commute (i.e., $M_1 M_2 = M_2 M_1$)
- ▶ M_1 and M_2 are simultaneously diagonalizable.



System Model (cont.)

Per-step cost


$$c(x_t, u_t) = \sum_{i,j \in \mathcal{N}} [\mathbf{h}_q^{ij}(x_t^i)^\top Q(x_t^j) + \mathbf{h}_r^{ij}(u_t^i)^\top Q(u_t^j)]$$

where $H_q = [h_q^{ij}]$ and $H_r = [h_r^{ij}]$ are **symmetric matrices which have the same eigenvectors as M** .


Important special case

$$\triangleright H_q = \sum_{k=0}^{K_q} q_k M^k \text{ and } H_r = \sum_{k=0}^{K_r} r_k M^k.$$

- \triangleright Captures the intuition that the per-step cost respects the graph structure.
- \triangleright Example: $H_q = q_0 I + q_1 M + q_2 M^2$ means that there is a cost coupling between the one- and two-hop neighbors.



An example to illustrate that **nodes** are not exchangeable



A graph G

An example to illustrate that nodes are not exchangeable

Dynamical coupling


- Nodes are not exchangeable

$$x_t^{\mathcal{G},1} = 2x_t^2 + 1x_t^4,$$

$$x_t^{\mathcal{G},2} = 2x_t^1 + 2x_t^3,$$

$$x_t^{\mathcal{G},3} = 2x_t^2 + 1x_t^4,$$

$$x_t^{\mathcal{G},4} = 1x_t^1 + 1x_t^3.$$



A graph \mathcal{G}

An example to illustrate that nodes are not exchangeable

Dynamical coupling


- Nodes are not exchangeable

$$x_t^{g,1} = 2x_t^2 + 1x_t^4,$$

$$x_t^{g,2} = 2x_t^1 + 2x_t^3,$$

$$x_t^{g,3} = 2x_t^2 + 1x_t^4,$$

$$x_t^{g,4} = 1x_t^1 + 1x_t^3.$$



A graph \mathcal{G}

Cost coupling

- Nodes are not exchangeable

Suppose $H_q = q_0 I + q_1 M + q_2 M^2$.

An example to illustrate that nodes are not exchangeable

Dynamical coupling


- Nodes are not exchangeable

$$x_t^{g,1} = 2x_t^2 + 1x_t^4,$$

$$x_t^{g,2} = 2x_t^1 + 2x_t^3,$$

$$x_t^{g,3} = 2x_t^2 + 1x_t^4,$$

$$x_t^{g,4} = 1x_t^1 + 1x_t^3.$$




A graph \mathcal{G}

Cost coupling

- Nodes are not exchangeable

Suppose $H_q = q_0 I + q_1 M + q_2 M^2$.



Two-hop neighborhood


An example to illustrate that nodes are not exchangeable

Dynamical coupling

- Nodes are not exchangeable

$$x_t^{g,1} = 2x_t^2 + 1x_t^4, \quad x_t^{g,2} = 2x_t^1 + 2x_t^3,$$

$$x_t^{g,3} = 2x_t^2 + 1x_t^4, \quad x_t^{g,4} = 1x_t^1 + 1x_t^3.$$




A graph \mathcal{G}

Cost coupling

- Nodes are not exchangeable

Suppose $H_q = q_0 I + q_1 M + q_2 M^2$. Then

$$H_q = \begin{bmatrix} q_0 + 5q_2 & 2q_1 & 5q_2 & q_1 \\ 2q_1 & q_0 + 8q_2 & 2q_1 & 4q_2 \\ 5q_2 & 2q_1 & q_0 + 5q_2 & q_1 \\ q_1 & 4q_2 & q_1 & q_0 + 2q_2 \end{bmatrix}$$




Two-hop neighborhood

Model generalizes mean-field control model

Special case

- ▶ Consider $M = \frac{1}{n} \mathbb{1}_{n \times n}$ and $H_q = H_r = \frac{1}{n} I + \frac{\kappa}{n} M$.
- ▶ Network-field $\frac{1}{n} \sum_{j \in N} x_t^j =: \bar{x}_t$ is the (empirical) mean-field.



Model generalizes mean-field control model

Special case


- ▶ Consider $M = \frac{1}{n} \mathbb{1}_{n \times n}$ and $H_q = H_r = \frac{1}{n} I + \frac{\kappa}{n} M$.
- ▶ Network-field $\frac{1}{n} \sum_{j \in N} x_t^j =: \bar{x}_t$ is the (empirical) mean-field.

Dynamics

$$x_{t+1}^i = Ax_t^i + Bu_t^i + D\bar{x}_t + E\bar{u}_t + w_t^i.$$

Per-step cost


$$\begin{aligned} c(x_t, u_t) = & (1 + \kappa) [\bar{x}_t^\top Q \bar{x}_t + \bar{u}_t^\top R \bar{u}_t] \\ & + \frac{1}{n} \sum_{i \in N} [(x_t^i - \bar{x}_t)^\top Q (x_t^i - \bar{x}_t) + (u_t^i - \bar{u}_t)^\top Q (u_t^i - \bar{u}_t)]. \end{aligned}$$



Problem formulation

Summary of the model


- ▶ **Dynamics:** $x_{t+1}^i = Ax_t^i + Bu_t^i + D \sum_{j \in N} m^{ij} x_t^j + E \sum_{j \in N} m^{ij} u_t^j + w_t^i$
- ▶ **Per-step cost:** $c(x_t, u_t) = \sum_{i,j \in N} [h_q^{ij} (x_t^i)^\top Q(x_t^j) + h_r^{ij} (u_t^i)^\top Q(u_t^j)]$
- ▶ **Network structure:** M , H_q , and H_r have the same eigenvectors.



Problem formulation

Summary of the model

- ▶ **Dynamics:** $x_{t+1}^i = Ax_t^i + Bu_t^i + D \sum_{j \in N} m^{ij} x_t^j + E \sum_{j \in N} m^{ij} u_t^j + w_t^i$
- ▶ **Per-step cost:** $c(x_t, u_t) = \sum_{i,j \in N} [h_q^{ij} (x_t^i)^\top Q(x_t^j) + h_r^{ij} (u_t^i)^\top Q(u_t^j)]$
- ▶ **Network structure:** M , H_q , and H_r have the same eigenvectors.



Objective


Choose a policy $\pi: (x_t^1, \dots, x_t^n) \rightarrow (u_t^1, \dots, u_t^n)$ to minimize:

$$\limsup_{T \rightarrow \infty} \frac{1}{T} \mathbb{E}^\pi \left[\sum_{t=1}^T c(x_t, u_t) \right]$$

Problem formulation

Summary of the model

- ▶ **Dynamics:** $x_{t+1}^i = Ax_t^i + Bu_t^i + D \sum_{j \in N} m^{ij} x_t^j + E \sum_{j \in N} m^{ij} u_t^j + w_t^i$
- ▶ **Per-step cost:** $c(x_t, u_t) = \sum_{i,j \in N} [h_q^{ij} (x_t^i)^\top Q(x_t^j) + h_r^{ij} (u_t^i)^\top Q(u_t^j)]$
- ▶ **Network structure:** M , H_q , and H_r have the same eigenvectors.




Objective

Choose a policy $\pi: (x_t^1, \dots, x_t^n) \rightarrow (u_t^1, \dots, u_t^n)$ to minimize:

$$\limsup_{T \rightarrow \infty} \frac{1}{T} \mathbb{E}^\pi \left[\sum_{t=1}^T c(x_t, u_t) \right]$$


- ▶ Standard soln requires solving $nd_x \times nd_x$ Riccati Eq.
- ▶ Complexity scales $\mathcal{O}(n^3 d_x^3)$.

Outline




System Model

- ▶ Network-coupled subsystems
- ▶ Agents interacting over a graph
- ▶ Coupled dynamics
- ▶ Coupled costs



Planning solution

- ▶ Spectral factorization
of dynamics and cost
- ▶ Decoupled Riccati equations



Learning solution

- ▶ Spectral factorization of learning
- ▶ Numerical experiments

Our result: Develop a decomposition which computes the optimal policy by solving at most n Riccati eqns of dimension $d_x \times d_x$.

- ▶ co-author: Shuang Gao
- ▶ paper: TCNS 2022

Spectral decomposition

Spectral decomposition of coupling matrices

$$\mathcal{M} = \sum_{\ell=1}^L \lambda^{\ell} \mathbf{v}^{\ell} (\mathbf{v}^{\ell})^{\top},$$

Spectral decomposition

Spectral decomposition of coupling matrices

$$M = \sum_{\ell=1}^L \lambda^\ell \mathbf{v}^\ell (\mathbf{v}^\ell)^\top, \quad H_q = q_0 I + q_1 \sum_{\ell=1}^L \lambda_q^\ell \mathbf{v}^\ell (\mathbf{v}^\ell)^\top, \quad H_r = r_0 I + r_1 \sum_{\ell=1}^L \lambda_r^\ell \mathbf{v}^\ell (\mathbf{v}^\ell)^\top$$

Spectral decomposition

Spectral decomposition of coupling matrices

$$M = \sum_{\ell=1}^L \lambda^\ell \mathbf{v}^\ell (\mathbf{v}^\ell)^\top, \quad H_q = q_0 I + q_1 \sum_{\ell=1}^L \lambda_q^\ell \mathbf{v}^\ell (\mathbf{v}^\ell)^\top, \quad H_r = r_0 I + r_1 \sum_{\ell=1}^L \lambda_r^\ell \mathbf{v}^\ell (\mathbf{v}^\ell)^\top$$

Spectral decomposition of dynamics

At each node $i \in [n]$:

► For each $\ell \in [L]$, define **eigenstates**, **eigencontrols**, and **eigennoise** as

$$\mathbf{x}_t^{\ell,i} = \mathbf{x}_t^i \mathbf{v}^\ell (\mathbf{v}^\ell)^\top, \quad \mathbf{u}_t^{\ell,i} = \mathbf{u}_t^i \mathbf{v}^\ell (\mathbf{v}^\ell)^\top, \quad \text{and} \quad \mathbf{w}_t^{\ell,i} = \mathbf{w}_t^i \mathbf{v}^\ell (\mathbf{v}^\ell)^\top.$$

Spectral decomposition

Spectral decomposition of coupling matrices

$$M = \sum_{\ell=1}^L \lambda^\ell \mathbf{v}^\ell (\mathbf{v}^\ell)^\top, \quad H_q = q_0 I + q_1 \sum_{\ell=1}^L \lambda_q^\ell \mathbf{v}^\ell (\mathbf{v}^\ell)^\top, \quad H_r = r_0 I + r_1 \sum_{\ell=1}^L \lambda_r^\ell \mathbf{v}^\ell (\mathbf{v}^\ell)^\top$$

Spectral decomposition of dynamics

At each node $i \in [n]$:

► For each $\ell \in [L]$, define **eigenstates**, **eigencontrols**, and **eigennoise** as

$$\mathbf{x}_t^{\ell,i} = \mathbf{x}_t^i \mathbf{v}^\ell (\mathbf{v}^\ell)^\top, \quad \mathbf{u}_t^{\ell,i} = \mathbf{u}_t^i \mathbf{v}^\ell (\mathbf{v}^\ell)^\top, \quad \text{and} \quad \mathbf{w}_t^{\ell,i} = \mathbf{w}_t^i \mathbf{v}^\ell (\mathbf{v}^\ell)^\top.$$

► Define **auxiliary state**, **auxiliary control**, **auxiliary noise** as

$$\check{\mathbf{x}}_t^i = \mathbf{x}_t^i - \sum_{\ell=1}^L \mathbf{x}_t^{\ell,i}, \quad \check{\mathbf{u}}_t^i = \mathbf{u}_t^i - \sum_{\ell=1}^L \mathbf{u}_t^{\ell,i}, \quad \text{and} \quad \check{\mathbf{w}}_t^i = \mathbf{w}_t^i - \sum_{\ell=1}^L \mathbf{w}_t^{\ell,i}.$$

Implication of Spectral Decomposition

Noise-coupled
dynamics

$$x_{t+1}^{\ell,i} = (A + \lambda^\ell D) x_t^{\ell,i} + (B + \lambda^\ell E) u_t^{\ell,i} + w_t^{\ell,i}$$

and $\check{x}_{t+1}^i = A\check{x}_t^i + B\check{u}_t^i + \check{w}_t^i$

Implication of Spectral Decomposition

Noise-coupled
dynamics

$$x_{t+1}^{\ell,i} = (A + \lambda^\ell D) x_t^{\ell,i} + (B + \lambda^\ell E) u_t^{\ell,i} + w_t^{\ell,i}$$

and $\check{x}_{t+1}^i = A\check{x}_t^i + B\check{u}_t^i + \check{w}_t^i$

Decoupled cost

$$c(x_t, u_t) = \sum_{i \in \mathcal{N}} \left[q_0 \check{c}(\check{x}_t^i, \check{u}_t^i) + \sum_{\ell=1}^L q^\ell c^\ell(x_t^{\ell,i}, u_t^{\ell,i}) \right]$$

where $q^\ell = q_0 + q_1 \lambda_q^\ell$, $r^\ell = r_0 + r_1 \lambda_r^\ell$, and

$$\check{c}(\check{x}_t^i, \check{u}_t^i) = (\check{x}_t^i)^\top Q \check{x}_t^i + \frac{r_0}{q_0} (\check{u}_t^i)^\top R \check{u}_t^i$$

$$c^\ell(x_t^{\ell,i}, u_t^{\ell,i}) = (x_t^{\ell,i})^\top Q x_t^{\ell,i} + \frac{r^\ell}{q^\ell} (u_t^{\ell,i})^\top R u_t^{\ell,i}.$$

Implication of Spectral Decomposition

Eigen-system (ℓ, i) with $\ell \in [L], i \in [n]$

- ▶ State $x_t^{\ell, i}$. Control $u_t^{\ell, i}$.
- ▶ Dynamics: $x_{t+1}^{\ell, i} = (A + \lambda^\ell D)x_t^{\ell, i} + (B + \lambda^\ell E)u_t^{\ell, i} + w_t^{\ell, i}$
- ▶ Per-step cost: $c^\ell(x_t^{\ell, i}, u_t^{\ell, i})$.

Auxiliary system i with $i \in [n]$

- ▶ State \check{x}_t^i . Control \check{u}_t^i .
- ▶ Dynamics: $\check{x}_{t+1}^i = A\check{x}_t^i + B\check{u}_t^i + \check{w}_t^i$
- ▶ Per-step cost: $c^\ell(\check{x}_t^i, \check{u}_t^i)$.

Implication of Spectral Decomposition

Eigen-system (ℓ, i) with $\ell \in [L], i \in [n]$

- ▶ State $x_t^{\ell, i}$. Control $u_t^{\ell, i}$.
- ▶ Dynamics: $x_{t+1}^{\ell, i} = (A + \lambda^\ell D)x_t^{\ell, i} + (B + \lambda^\ell E)u_t^{\ell, i} + w_t^{\ell, i}$
- ▶ Per-step cost: $c^\ell(x_t^{\ell, i}, u_t^{\ell, i})$.

Auxiliary system i with $i \in [n]$

- ▶ State \tilde{x}_t^i . Control \tilde{u}_t^i .
- ▶ Dynamics: $\tilde{x}_{t+1}^i = A\tilde{x}_t^i + B\tilde{u}_t^i + \tilde{w}_t^i$
- ▶ Per-step cost: $c^\ell(\tilde{x}_t^i, \tilde{u}_t^i)$.

Only coupled through the noise in the dynamics

Implication of Spectral Decomposition

Eigen-system (ℓ, i) with $\ell \in [L], i \in [n]$

- ▶ State $x_t^{\ell, i}$. Control $u_t^{\ell, i}$.
- ▶ Dynamics: $x_{t+1}^{\ell, i} = (A + \lambda^\ell D)x_t^{\ell, i} + (B + \lambda^\ell E)u_t^{\ell, i} + w_t^{\ell, i}$
- ▶ Per-step cost: $c^\ell(x_t^{\ell, i}, u_t^{\ell, i})$.

Auxiliary system i with $i \in [n]$

- ▶ State \check{x}_t^i . Control \check{u}_t^i .
- ▶ Dynamics: $\check{x}_{t+1}^i = A\check{x}_t^i + B\check{u}_t^i + \check{w}_t^i$
- ▶ Per-step cost: $c^\ell(\check{x}_t^i, \check{u}_t^i)$.

Certainty equivalence: Optimal policy of stochastic LQ system is same as that of deterministic LQ system.

The deterministic system has **decoupled dynamics and cost!**

Only coupled through the noise in the dynamics

Main result

Under standard assumptions, the optimal control action is given by

$$u_t^i = \check{u}_t^i + \sum_{\ell=1}^L u_t^{\ell,i} = \check{G} \check{x}_t^i + \sum_{\ell=1}^L G^\ell x_t^{\ell,i}$$

where

$$\check{G} = \text{Gain} \left(A, B, Q, \frac{r_0}{q_0} R \right)$$

$$G^\ell = \text{Gain} \left(A + \lambda^\ell D, B + \lambda^\ell E, Q, \frac{r^\ell}{q^\ell} R \right), \quad \ell \in [L]$$

Main result

Under standard assumptions, the optimal control action is given by

$$u_t^i = \check{u}_t^i + \sum_{\ell=1}^L u_t^{\ell,i} = \check{G} \check{x}_t^i + \sum_{\ell=1}^L G^\ell x_t^{\ell,i}$$


where

$$\check{G} = \text{Gain}\left(A, B, Q, \frac{r_0}{q_0} R\right)$$

$$G^\ell = \text{Gain}\left(A + \lambda^\ell D, B + \lambda^\ell E, Q, \frac{r^\ell}{q^\ell} R\right), \quad \ell \in [L]$$


- ▶ The gains $\check{G}, \{G^\ell\}_{\ell=1}^L$ are the same at all subsystems!
- ▶ Requires solving **(L + 1) Riccati Eqn** of dimension $d_x \times d_x$.
- ▶ Complexity scales **$\mathcal{O}(L d_x^3)$** (cf. $\mathcal{O}(n^3 d_x^3)$ for naive solution).

Outline




System Model

- ▶ Network-coupled subsystems
- ▶ Agents interacting over a graph
- ▶ Coupled dynamics
- ▶ Coupled costs



Planning solution

- ▶ Spectral factorization
of dynamics and cost
- ▶ Decoupled Riccati equations



Learning solution

- ▶ Spectral factorization of learning
- ▶ Numerical experiments

Review: LQ regulation with unknown/uncertain dynamics

Modeling
uncertainty

Model $\theta_* = [A_*, B_*] \in \Theta$ (uncertain set)

Review: LQ regulation with unknown/uncertain dynamics

Modeling uncertainty

Model $\theta_* = [A_*, B_*] \in \Theta$ (uncertain set)

Worst-case design

- ▶ Assume that nature is adversarial
- ▶ Choose a policy which minimizes worst case performance
- ▶ **Zero-sum game or robust control**

Review: LQ regulation with unknown/uncertain dynamics

Modeling uncertainty

Model $\theta_* = [A_*, B_*] \in \Theta$ (uncertain set)

Worst-case design

- ▶ Assume that nature is adversarial
- ▶ Choose a policy which minimizes worst case performance
- ▶ **Zero-sum game or robust control**

Learning solution

- ▶ Design an adaptive policy which asymptotically converges to the optimal policy of the true (unknown) model.
- ▶ **Reinforcement learning or adaptive control**

Review: LQ regulation with unknown/uncertain dynamics

Modeling uncertainty

Model $\theta_* = [A_*, B_*] \in \Theta$ (uncertain set)

Worst-case design

- ▶ Assume that nature is adversarial
- ▶ Choose a policy which minimizes worst case performance
- ▶ **Zero-sum game or robust control**

Learning solution

- ▶ Design an adaptive policy which asymptotically converges to the optimal policy of the true (unknown) model.
- ▶ **Reinforcement learning or adaptive control**

Comparing learning algorithms

$$\text{Regret}(T) = \sum_{t=1}^T \left[\text{cost of learning algo}(t) - \text{cost of clairvoyant agent}(t) \right]$$

- ▶ Large literature. Various classes of algos with different regret guarantees.

Review: Regret for learning in LQ regulation

Bounds on Regret

- ▶ **Lower bound:** No algorithm can do better than $\tilde{\Omega}(d_x^{0.5} d_u \sqrt{T})$.
- ▶ **Upper bound:** Various classes of algorithms achieve $\tilde{O}(d_x^{0.5} (d_x + d_u) \sqrt{T})$.
 - ▶ Certainty equivalence; Optimisim in the face of uncertainty; Thompson sampling

Review: Regret for learning in LQ regulation

Bounds on Regret

- ▶ **Lower bound:** No algorithm can do better than $\tilde{\Omega}(d_x^{0.5} d_u \sqrt{T})$.
- ▶ **Upper bound:** Various classes of algorithms achieve $\tilde{O}(d_x^{0.5} (d_x + d_u) \sqrt{T})$.
 - ▶ Certainty equivalence; Optimisim in the face of uncertainty; Thompson sampling

Challenge with learning in networks

- ▶ Effective dimensions are nd_x and nd_u
- ▶ Directly using existing algos gives regret of $\tilde{O}(n^{1.5} d_x^{0.5} (d_x + d_u) \sqrt{T})$.
- ▶ **Normalized** regret per agent is $\tilde{O}(n^{0.5} d_x^{0.5} (d_x + d_u) \sqrt{T})$.

Review: Regret for learning in LQ regulation

Bounds on Regret

- ▶ **Lower bound:** No algorithm can do better than $\tilde{\Omega}(d_x^{0.5} d_u \sqrt{T})$.
- ▶ **Upper bound:** Various classes of algorithms achieve $\tilde{O}(d_x^{0.5} (d_x + d_u) \sqrt{T})$.
 - ▶ Certainty equivalence; Optimisim in the face of uncertainty; Thompson sampling

Challenge with learning in networks

- ▶ Effective dimensions are nd_x and nd_u
- ▶ Directly using existing algos gives regret of $\tilde{O}(n^{1.5} d_x^{0.5} (d_x + d_u) \sqrt{T})$.
- ▶ **Normalized** regret per agent is $\tilde{O}(n^{0.5} d_x^{0.5} (d_x + d_u) \sqrt{T})$.

Regret per agent grows with size of the network!

Our result: Develop a learning algorithm which exploits the structure of the network and has a per agent regret of $\tilde{O}\left((1 + \frac{1}{n}) d_x^{0.5} (d_x + d_u) \sqrt{T}\right)$.

- ▶ [co-author](#): Sagar Sudhakara, Ashutosh Nayyar, Yi Ouyang
- ▶ [paper](#): TCNS 2023

Learning model

Problem setting

- ▷ **Known:** Network (M, H_q, H_r) . Cost (Q, R) .
- ▷ **Unknown:** Dynamics (A, B, C, D) .

Learning model

Problem setting

- ▶ **Known:** Network (M, H_q, H_r) . Cost (Q, R) .
- ▶ **Unknown:** Dynamics (A, B, C, D) .

Modeling assumptions

- ▶ $\check{\theta}_* = [A_*, B_*] \in \check{\Theta}$
- ▶ $\theta^\ell = [A_* + \lambda^\ell D_*, B_* + \lambda^\ell E_*] \in \Theta^\ell$
- ▶ Bayesian prior on $\check{\Theta}$ and $\{\Theta^\ell\}_{\ell=1}^L$.

Learning model

Problem setting

- ▶ **Known:** Network (M, H_q, H_r) . Cost (Q, R) .
- ▶ **Unknown:** Dynamics (A, B, C, D) .

Modeling assumptions

- ▶ $\check{\theta}_* = [A_*, B_*] \in \check{\Theta}$
- ▶ $\theta^\ell = [A_* + \lambda^\ell D_*, B_* + \lambda^\ell E_*] \in \Theta^\ell$
- ▶ Bayesian prior on $\check{\Theta}$ and $\{\Theta^\ell\}_{\ell=1}^L$.

Implication of Spectral Decomposition

$$\text{Recall: } c(x_t, u_t) = \sum_{i \in N} \left[\mathbf{q}_0 \check{c}(\check{x}_t^i, \check{u}_t^i) + \sum_{\ell=1}^L \mathbf{q}^\ell c^\ell(x_t^{\ell,i}, u_t^{\ell,i}) \right]$$

Thus, for any policy π ,

$$J(\pi; \theta_*) = \sum_{i \in N} \left[\mathbf{q}_0 \check{J}^i(\pi; \check{\theta}_*) + \sum_{\ell=1}^L \mathbf{q}^\ell J^{\ell,i}(\pi; \theta_*^\ell) \right].$$

Key idea for learning

Separately learn $\{\theta^\ell\}_{\ell=1}^L$ and $\check{\theta}$

- ▶ For learning θ_*^ℓ , select an agent i_o^ℓ such that $v^{\ell, i_o^\ell} \neq 0$.
- ▶ Learn $G^\ell(\theta_*^\ell)$ using $\{x_t^{\ell, i_o^\ell}, u_t^{\ell, i_o^\ell}\}_{t \geq 1}$.

Key idea for learning

Separately learn $\{\theta^\ell\}_{\ell=1}^L$ and $\check{\theta}$

- ▶ For learning θ_*^ℓ , select an agent i_o^ℓ such that $v^{\ell, i_o^\ell} \neq 0$.
- ▶ Learn $G^\ell(\theta_*^\ell)$ using $\{x_t^{\ell, i_o^\ell}, u_t^{\ell, i_o^\ell}\}_{t \geq 1}$.
- ▶ At time t , select agent j_{t-1} with the “most informative obs”.
- ▶ Learn $\check{G}(\check{\theta}_*)$ using $\{\check{x}_t^{j_t}, \check{u}_t^{j_t}, \check{x}_{t+1}^{j_t}\}_{t \geq 1}$.

Key idea for learning

Separately learn $\{\theta^\ell\}_{\ell=1}^L$ and $\check{\theta}$

- ▶ For learning θ_*^ℓ , select an agent i_o^ℓ such that $v^{\ell, i_o^\ell} \neq 0$.
- ▶ Learn $G^\ell(\theta_*^\ell)$ using $\{x_t^{\ell, i_o^\ell}, u_t^{\ell, i_o^\ell}\}_{t \geq 1}$.
- ▶ At time t , select agent j_{t-1} with the “most informative obs”.
- ▶ Learn $\check{G}(\check{\theta}_*)$ using $\{\check{x}_t^{j_t}, \check{u}_t^{j_t}, \check{x}_{t+1}^{j_t}\}_{t \geq 1}$.

- ▶ Use variant of Thompson sampling to learn each component
- ▶ The high-level idea also applies to other learning algos

Key idea for learning

Separately learn $\{\theta^\ell\}_{\ell=1}^L$ and $\check{\theta}$

- ▶ For learning θ_*^ℓ , select an agent i_o^ℓ such that $v^{\ell, i_o^\ell} \neq 0$.
- ▶ Learn $G^\ell(\theta_*^\ell)$ using $\{x_t^{\ell, i_o^\ell}, u_t^{\ell, i_o^\ell}\}_{t \geq 1}$.
- ▶ At time t , select agent j_{t-1} with the “most informative obs”.
- ▶ Learn $\check{G}(\check{\theta}_*)$ using $\{\check{x}_t^{j_t}, \check{u}_t^{j_t}, \check{x}_{t+1}^{j_t}\}_{t \geq 1}$.

- ▶ Use variant of Thompson sampling to learn each component
- ▶ The high-level idea also applies to other learning algos

Implication of Spectral Decomposition

$$\text{Since } J(\pi; \theta_*) = \sum_{i \in N} \left[\mathbf{q}_0 \check{J}^i(\pi; \check{\theta}_*) + \sum_{\ell=1}^L \mathbf{q}^\ell J^{\ell, i}(\pi; \theta_*^\ell) \right].$$

Thus, regret also decomposes as

$$R(T) = \sum_{i \in N} \left[\mathbf{q}_0 \check{R}^i(T) + \sum_{\ell=1}^L \mathbf{q}^\ell R^{\ell, i}(T) \right].$$

Bounding regret

Bounding $R^{\ell, i}(T)$

- ▶ Since agent i_o^ℓ is learning in the standard manner, we have

$$R^{\ell, i_o^\ell}(T) = \tilde{O}(\mathbf{W}^{\ell, i_o^\ell} d_x^{0.5} (d_x + d_u) \sqrt{T}).$$

- ▶ We show that for other agents

$$R^{\ell, i}(T) = \left(\frac{v^{\ell, i}}{v^{\ell, i_o^\ell}} \right)^2 R^{\ell, i_o^\ell}(T) = \tilde{O}(\mathbf{W}^{\ell, i} d_x^{0.5} (d_x + d_u) \sqrt{T}).$$

Bounding regret

Bounding $R^{\ell, i}(T)$

- ▶ Since agent i_o^ℓ is learning in the standard manner, we have

$$R^{\ell, i_o^\ell}(T) = \tilde{O}(\mathbf{W}^{\ell, i_o^\ell} d_x^{0.5} (d_x + d_u) \sqrt{T}).$$

- ▶ We show that for other agents

$$R^{\ell, i}(T) = \left(\frac{v^{\ell, i}}{v^{\ell, i_o^\ell}} \right)^2 R^{\ell, i_o^\ell}(T) = \tilde{O}(\mathbf{W}^{\ell, i} d_x^{0.5} (d_x + d_u) \sqrt{T}).$$

Bounding $\check{R}^i(T)$

- ▶ Need to bound regret from first principles.
- ▶ Using the most informative observation allows us to bound the regret of auxiliary systems at all nodes.
- ▶ Show that $\check{R}^i(T) = \tilde{O}(\check{\mathbf{W}}^i d_x^{0.5} (d_x + d_u) \sqrt{T})$.

Bounding regret

▶ Since agent i_0^ℓ is learning in the standard manner, we have

$$R(i_0^\ell(T)) = \tilde{O}(\alpha^\ell d_x^{0.5} (d_x + d_u) \sqrt{T})$$

Bounding R^ℓ

Overall Regret Bound

▶ Combining these, we have

$$R(T) = \tilde{O}(\alpha^{\mathcal{G}} d_x^{0.5} (d_x + d_u) \sqrt{T}),$$

where $\alpha^{\mathcal{G}} = \sum_{\ell=1}^L q^\ell + q_0(n-L).$

▶ Regret per agent is proportional to

$$\alpha^{\mathcal{G}}/n = \mathcal{O}\left(1 + \frac{L}{n}\right).$$

Thus, regret per agent reduces with the size of the network!

Bounding \tilde{R}


the

Some examples

Mean-field systems

Choice of parameters

▷ $M = \frac{1}{n} \mathbb{1}_{n \times n}$ and $H_q = H_r = \frac{1}{n} I + \frac{\kappa}{n} M$.



Mean-field systems


Choice of parameters

▷ $M = \frac{1}{n} \mathbb{1}_{n \times n}$ and $H_q = H_r = \frac{1}{n} I + \frac{\kappa}{n} M$.

Scaling of regret

▷ $q^1 = r^1 = (1 + \kappa)/n$. Thus, (normalized) $\alpha^g = \left(1 + \frac{\kappa}{n}\right)$.

▷ Regret per-agent goes down as the network becomes larger (mean-field effect).






Mean-field systems

Choice of parameters

▷ $M = \frac{1}{n} \mathbb{1}_{n \times n}$ and $H_q = H_r = \frac{1}{n} I + \frac{\kappa}{n} M$.

Scaling of regret



▷ $q^1 = r^1 = (1 + \kappa)/n$. Thus, (normalized) $\alpha^g = \left(1 + \frac{\kappa}{n}\right)$.



A general low-rank network

Choice of parameters

► $M = M^\circ \otimes \frac{1}{n} \mathbb{1}_{n \times n}$, $H_q = (I - M)^2$, and $H_r = I$.





A general low-rank network

Choice of parameters

- ▶ $M = M^\circ \otimes \frac{1}{n} \mathbb{1}_{n \times n}$, $H_q = (I - M)^2$, and $H_r = I$.

Scaling of regret



- ▶ $\lambda^\ell = \pm \sqrt{2(a^2 + b^2)}$, $q^\ell = (1 - \lambda^\ell)^2$, $r^\ell = 1$. Thus, (unnormalized) $\alpha^{\mathcal{G}} = 4n + 4(a^2 + b^2)$.
- ▶ Regret per-agent goes down as the network becomes larger (network-field effect).



A general low-rank network



Choice of parameters

► $M = M^\circ \otimes \frac{1}{n} \mathbb{1}_{n \times n}$, $H_q = (I - M)^2$, and $H_r = I$.



Scaling of regret

► $\lambda^\ell = \pm \sqrt{2(a^2 + b^2)}$, $q^\ell = (1 - \lambda^\ell)^2$, $r^\ell = 1$. Thus, (unnormalized) $\alpha^g = 4n + 4(a^2 + b^2)$.



Conclusion

Presented a spectral decomposition method for network-coupled subsystems which leads to scalable planning and learning

Conclusion

Presented a spectral decomposition method for network-coupled subsystems which leads to scalable planning and learning

Planning solution

- ▶ Solve $(L + 1)$ Riccati eqns of dims $d_x \times d_x$.

Learning solution

- ▶ Regret per agent $\tilde{O}((1 + \frac{1}{n})\sqrt{T})$

Conclusion

Presented a spectral decomposition method for network-coupled subsystems which leads to scalable planning and learning

Planning solution

- ▶ Solve $(L + 1)$ Riccati eqns of dims $d_x \times d_x$.

Learning solution

- ▶ Regret per agent $\tilde{O}((1 + \frac{1}{n})\sqrt{T})$

Future Directions

- ▶ Multiple types of agents, approximate symmetry, ...
- ▶ Large networks, graphon limits?
- ▶ Other types of scalable network structures?

- ▶ email: aditya.mahajan@mcgill.ca
- ▶ web: <http://cim.mcgill.ca/~adityam>

Thank you

Funding

- ▶ NSERC Discovery
- ▶ DND IDEaS Network

References

- ▶ planning: TCNS 2022
- ▶ learning: TCNS 2023