

Non-classical information structures

Examples of sequential decomposition

Aditya Mahajan

Yale University

Thanks to

Ashutosh Nayyar and Demosthenis Teneketzis,
University of Michigan

Information Structures

- Information structures :

Collection of data available to each agent to make a decision

- Classical information structure (Centralized Systems)

An agent knows the data available
to all agents that acted earlier

- Non-classical information structure (Decentralized Systems)



Centralized stochastic control – The two models

- Models

- ▷ LQG
- ▷ Control of Markov Chains

- Structural results

- ▷ (LQG) Affine control laws are optimal
- ▷ (MC) Controls based on belief state are optimal

- Sequential decomposition

Break a one-shot optimization into a sequence of nested optimizations

- ▷ Dynamic programming



Outline

- Sequential decomposition in centralized stochastic control
 - ▷ MDP (Markov decision processes)
 - ▷ POMDP (Partially observable Markov decision processes)
- Sequential decomposition in decentralized stochastic control
 - ▷ Common observations
 - ▷ No common observation
- Examples
 - ▷ Decentralized detection
 - ▷ Queueing theory
 - ▷ Robotics



Centralized stochastic control – MDP

- Model

State Update : $X(t+1) = f_t(X(t), U(t), W(t))$

Observation : $Y(t) = X(t)$

Control : $U(t) = g_t(X(1:t), U(1:t-1))$

Objective : $E^G \left\{ \sum_{t=1}^T c_t(X(t), U(t)) \right\}$

- Structural Result

Control : $U(t) = g_t(X(t))$

- Sequential Decomposition

$$V_t(x(t)) = \min_{u(t)} [c_t(x(t), u(t)) + E\{V_{t+1}(X(t+1)) \mid x(t), u(t)\}]$$



Centralized stochastic control – POMDP

○ Model

State Update : $X(t+1) = f_t(X(t), U(t), W(t))$

Observation : $Y(t) = h_t(X(t), N(t))$

Control : $U(t) = g_t(Y(1:t), U(1:t-1))$

Objective : $E^G \left\{ \sum_{t=1}^T c_t(X(t), U(t)) \right\}$

○ Structural Result

Control : $U(t) = g_t(B(t))$

where $B(t)(x) = \Pr(X(t) = x | Y(1:t), U(1:t-1))$

○ Sequential Decomposition

$$V_t(b(t)) = \min_{u(t)} [E\{c_t(x(t), u(t)) + V_{t+1}(B(t+1) | b(t), u(t))\}]$$



Decentralized Systems

Decentralized stochastic control

- Cost of decentralization
 - ▷ Triple objective of control: control, estimation, and **communication**
- No easy way to find structural results
 - ▷ Can consider the problem from the p.o.v. of one agent and fix the strategies of all other agents
 - ▷ Works for two agent problems
 - ▷ May not work in general and may not give **compact** structural results
- No general way to obtain sequential decomposition (until now)
 - ▷ Coupled dynamic programs can give locally optimal strategies



Decentralized stochastic control

- Our approach (Mahajan, Nayyar, Teneketzis 2008)

Find sequential decomposition of specific non-trivial instances.

- Two models

- ▷ Model A — Common observations
- ▷ Model B — No common observations

- Cost of decentralization

- ▷ Triple objective of control: control, estimation, and **communication**
- ▷ Each step of sequential decomposition is a functional optimization problem



Model A – Common observations

- State Update:

$$X(t+1) = f_t(X(t), U(t), W(t)) \quad \text{where } U(t) = [U_1(t), \dots, U_n(t)]$$

- Common Observations: $Z(t) = c_t(X(t), Q(t))$
- Private Observations: $Y_i(t) = h_{i,t}(X(t), N(t))$
- Control: $U_i(t) = g_{i,t}(Z(1:t-1), Y_i(t), M_i(t-1))$
- Memory: $M_i(t) = l_{i,t}(Z(1:t-1), Y_i(t), M_i(t-1))$
- Objective: $E^{G,L} \left\{ \sum_{t=1}^T c_t(X(t), U(t)) \right\}$



Main Idea:
Think in term of
common agent

Common agent

- Partial functions

$$g_{i,t} : \mathcal{Z}^{t-1} \times \mathcal{Y}_i \times \mathcal{M}_i \rightarrow \mathcal{U}_i$$

$$g_{i,t} : \mathcal{Z}^{t-1} \rightarrow (\mathcal{Y}_i \times \mathcal{M}_i \rightarrow \mathcal{U}_i)$$

$$l_{i,t} : \mathcal{Z}^{t-1} \times \mathcal{Y}_i \times \mathcal{M}_i \rightarrow \mathcal{M}_i$$

$$l_{i,t} : \mathcal{Z}^{t-1} \rightarrow (\mathcal{Y}_i \times \mathcal{M}_i \rightarrow \mathcal{M}_i)$$

- Common agent decides partial functions

$$\mathcal{U}_i(t) = \hat{g}_{i,t}(\mathcal{Y}_i(t), \mathcal{M}_i(t-1)) \quad \text{where} \quad \hat{g}_{i,t} = \gamma_{i,t}(z(1:t))$$

$$\mathcal{M}_i(t) = \hat{l}_{i,t}(\mathcal{Y}_i(t), \mathcal{M}_i(t-1)) \quad \text{where} \quad \hat{l}_{i,t} = \lambda_{i,t}(z(1:t))$$



Equivalent model

- Augmented State : $(X(t), Y(t), M(t-1))$

where $Y(t) = [Y_1(t), \dots, Y_n(t)]$ and same for $M(t)$.

- State Update

$$X(t+1) = f_t(X(t), U(t), W(t)) \quad Y_i(t) = h_{i,t}(X(t), N(t))$$

$$U_i(t) = \hat{g}_{i,t}(Y_i(t), M_i(t-1)), \quad M_i(t) = \hat{l}_{i,t}(Y_i(t), M_i(t-1))$$

- Observations: $Z(t) = c_t(X(t), U(t-1), Q(t))$

- Control $\hat{g}_{i,t} = \gamma_{i,t}(Z(1:t-1))$

$$\hat{l}_{i,t} = \lambda_{i,t}(Z(1:t-1))$$

- Objective $E^{\Gamma, \Lambda} \left\{ \sum_{t=1}^T c_t(X(t), Y(t), M(t-1), \hat{g}_{i,t}, \hat{l}_{i,t}) \right\}$



Equivalent model

Equivalent to a centralized problem!

- Information state

$$\pi(t)(z(1:t)) = \left[[\pi(t)(x, y, m)] \right]$$

where

$$\pi(t)(x, y, m) = \Pr(X(t) = x, Y(t) = y, M(t-1) = m \mid Z(1:t) = z(1:t))$$

- Sequential decomposition

$$V_t(\pi(t)) = \min_{\hat{g}(t), \hat{l}(t)} \left[E\{c_t(x(t), u(t)) + V_{t+1}(\Pi(t+1) \mid \pi(t), \hat{g}(t), \hat{l}(t))\} \right]$$



Model B – No common observations

- State Update:

$$X(t+1) = f_t(X(t), U(t), W(t)) \quad \text{where } U(t) = [U_1(t), \dots, U_n(t)]$$

- Observations: $Y_i(t) = h_{i,t}(X(t), N(t))$

- Control: $U_i(t) = g_{i,t}(Y_i(t), M_i(t-1))$

- Memory: $M_i(t) = g_{i,t}(Y_i(t), M_i(t-1))$

- Objective: $E^{G,L} \left\{ \sum_{t=1}^T c_t(X(t), U(t)) \right\}$



Main Idea:

Think in term of
system designer

Equivalent model

- Augmented State : $(X(t), Y(t), M(t-1))$

where $Y(t) = [Y_1(t), \dots, Y_n(t)]$ and same for $M(t)$.

- State Update

$$X(t+1) = f_t(X(t), U(t), W(t)) \quad Y_i(t) = h_{i,t}(X(t), N(t))$$

$$U_i(t) = g_{i,t}(Y_i(t), M_i(t-1)), \quad M_i(t) = l_{i,t}(Y_i(t), M_i(t-1))$$

- Observations: **Nothing**

- Control $g_{i,t} = \gamma_{i,t}(g(t-1), l(t-1))$

$$l_{i,t} = \lambda_{i,t}(g(t-1), l(t-1))$$

- Objective $E^{\Gamma, \Lambda} \left\{ \sum_{t=1}^T c_t(X(t), Y(t), M(t-1), g_{i,t}, l_{i,t}) \right\}$



Equivalent model

Equivalent to a centralized problem!

- Information state

$$\pi(t)(z(1:t)) = \left[\left[\pi(t)(z(1:t))(x, y, m) \right] \right]$$

where

$$\pi(t)(z(1:t))(x, y, m) = \Pr (X(t) = x, Y(t) = y, M(t-1) = m)$$

- Sequential decomposition

$$V_t(\pi(t)) = \min_{g(t), l(t)} \left[E \{ c_t(x(t), u(t)) + V_{t+1}(\Pi(t+1) \mid \pi(t), g(t), l(t)) \} \right]$$



Non-classical
information structures:
Sequential decomposition
is possible

Are these models useful – Special Cases

- k-step state sharing information structure
- k-step control sharing information structure
- k-step observation sharing information structure
- Witsenhausen's general sequential team model
- Witsenhausen's standard form



Are these models useful – Examples

- Decentralized Detection

- ▷ Decentralized Wald Problem
- ▷ Wald problem with 1-bit memory

- Control of Queues

- ▷ Multi-access broadcast
- ▷ Load balancing in queues

- Robotics

- ▷ Robot rendezvous with communications
- ▷ Decentralized task scheduling


- Toy problem

- ▷ Decentralized tiger problem



Decentralized Wald problem

- Decisions = { More meas., decide H_0 , decide H_1 }
- action = fn(all past observations)
- Each measurement by each sensor costs c
- Final cost $J(u_1, u_2, h) = \begin{cases} 0 & \text{if } u_1 = u_2 = h, \\ 1 & \text{if } u_1 \neq u_2 \\ k & \text{if } u_1 = u_2 \neq h \end{cases}$
- Minimize $E\{c\tau_1 + c\tau_2 + J(u_1, u_2, h)\}$



Decentralized Wald problem

- Studied by Teneketzis and Ho, 1985
- Obtained structural results:

Optimal to take action based on
likelihood ratio test (threshold rule)

$$\text{action} = \text{fn}(\text{posterior of } H_0)$$

- Obtained coupled dynamic programs to determine *locally optimally thresholds*.



Globally optimal thresholds – Use Model B

Structural results make the model equivalent to Model B

- Define: $\pi(t) = \Pr(H, U_1(t), U_2(t), \lambda_1(t), \lambda_2(t))$ or equivalently

$$\pi(t) = \left[\Pr(U_1(t), U_2(t), \lambda_1(t), \lambda_2(t) | H_0), \Pr(U_1(t), U_2(t), \lambda_1(t), \lambda_2(t) | H_1) \right]$$


- Sequential Decomposition

$$V_t(\pi(t)) = \min_{\alpha(t), \beta(t)} \left[E\{c_t(h, u_1(t), u_2(t)) + V_{t+1}(\Pi(t+1) | \pi(t), \alpha(t), \beta(t))\} \right]$$



Wald problem with 1 bit of memory

- Decisions = { More meas. and **update memory**, decide H_0 , decide H_1 }
- Each measurement by each sensor costs c
- (action, next memory) = fn(current observation, current memory)
- Final cost $J(u, h) = \begin{cases} 0 & \text{if } u = h, \\ L_0 & \text{if } u \neq h = H_0 \\ L_1 & \text{if } u \neq h = H_1 \end{cases}$
- Minimize $E\{c\tau + J(u, h)\}$



Wald problem with 1 bit of memory

- Studied by Sandell, 1974
- Presents a sequential decomposition attributed to Witsenhausen

Globally optimal design – Use Model B


- Define: $\pi(t) = \left[\Pr(M(t) | H_0), \Pr(M(t) | H_1) \right]$
- Sequential Decomposition

$$V_t(\pi(t)) = \min \left\{ \begin{array}{l} E\{\text{cost of stopping} \mid \pi(t), u = H_0\}, \\ E\{\text{cost of stopping} \mid \pi(t), u = H_1\}, \\ c + E\{V_{t+1}(\Pi(t+1)) \mid \pi(t), g(t)\} \end{array} \right\}$$



Multi-access broadcast

- Independent Bernoulli arrivals
- One user transmits \implies successful transmission
- Both users transmit \implies packet collision
- Channel feedback available to both users
- action = fn(all past queue sizes, all past feedback)
- Objective: Maximize throughput or minimize delay



Multi-access broadcast

- Studied by Hluchyj and Gallager, 1981
- Restricted attention to “window protocols” and symmetric arrival rates

Determining optimal design – Use Model A

- Structural results

action = fn(current queue size, all feedback)

- Define : $\pi(t) = \Pr(\text{queue size of user 1, queue size of user 2} \mid \text{all past feedback})$
- Sequential decomposition

$$V_t(\pi(t)) = \min_{g_1(t), g_2(t)} \left[E \left\{ c(x(t), u(t)) + V_{t+1}(\Pi(t+1)) \mid \pi(t), g_1(t), g_2(t) \right\} \right]$$



Multi-access broadcast


- For symmetric arrival rates and buffer size 1, the sequential decomposition can be used to prove that Hluchyj-Gallager strategy is optimal among all strategies.
- Structural results can be further simplified (Anastasopoulos, preprint)

$$\pi(t) = (\Pr(\text{queue size of user 1} \mid \text{all past feedback}), \\ \Pr(\text{queue size of user 2} \mid \text{all past feedback}))$$



Load balancing in Queues

- Independent Bernoulli arrivals at rate λ_i
- Geometric service times μ_i
- Each queue can look at the queue sizes of its neighbors and transfer packets to neighbors.
- Finite buffer.
- Waiting cost w , dropping cost d , switching cost c
- `action = fn(all past queue lengths of self and neighbors)`
- Objective: Minimize expected waiting plus switching cost



Load balancing in Queues

- Considered by Cogill, Rotkowitz, Roy, and Lall, 2004
- Assumed each queue uses only the current queue lengths to determine its action.

action = fn(current queue lengths of self and neighbors)

- Provided an approximate dynamic programming approach


Determining optimal design – Use model B

- Define: $\pi(t) = \Pr(\text{queue length of all users})$
- Sequential Decomposition

$$V_t(\pi(t)) = \min_{g(t)} [E\{c(x(t), u(t)) + V_{t+1}(\Pi(t+1)) \mid \pi(t), g(t)\}]$$




Robot rendezvous



Bernsein, Zibersein, Immerman, 2000



Robot task sharing



Conclusion

Finding structural results is an art

Obtaining a sequential decomposition
can be turned into a science